


Hybrid Volatility Surface Reconstruction

Integrating the SABR Model with TabPFN



Joël Khayat

Contents

1	Introduction	2
1.1	The SABR Volatility Model	2
1.2	The Machine Learning Paradigm: Introducing TabPFN	2
1.3	Bridging the Gap: Applying TabPFN to the SABR Model	3
1.3.1	The Greek Problem and the Need for a Hybrid Architecture:	3
2	SABR Engine Implementation (Steps 1 & 2)	4
2.1	Model Dynamics and Parameter Mapping (Step 1)	4
2.2	Hagan's Asymptotic Approximation and Limitations (Step 2)	4
2.3	Singularity Management and Autodiff Implementation	4
3	Synthetic Data Generation and Sobol Coverage (Step 3)	5
3.1	Parameter Space and Scaling	5
3.2	Dataset Analysis and Structural Integrity	5
4	Baseline Training: Standard TabPFN (Step 4)	7
4.1	Pre-processing and Scaling	7
4.2	Initial Performance and The "Gradient Challenge"	7
5	The Derivatives Issue: Baseline AI Evaluation	8
5.1	JAX as the Analytical Ground Truth	8
5.2	Comparative Analysis: The Gradient Challenge	8
6	Sobolev Loss Architecture: Learning the Surface Geometry (Step 5)	9
6.1	Mathematical Formulation	9
6.2	Implementation and Monitoring	9
7	Data Generation Strategy: Hybrid Mesh Refinement (Step 5)	10
7.1	Methodology: Sobol Foundation and Gaussian ATM Refinement	10
8	Architecture Search and Hyperparameter Optimization (Step 7)	11
8.1	Experimental Setup and Search Space	11
8.2	Optimal Configuration	12
9	Stacking (Step 8)	13
9.1	Stacking Architecture and Feature Augmentation	13
9.2	Sobolev Loss Regularization and JAX Ground Truth	13
10	Finetunning on Causal Graph	15
11	Results and Geometric Validation	15

1 Introduction

1.1 The SABR Volatility Model

Developed by Hagan, Kumar, Lesniewski, and Woodward in 2002, the **SABR** model is a two-factor stochastic volatility model. It models the forward price of an asset, F_t , and its instantaneous volatility, α_t , through the following system of stochastic differential equations (SDEs):

$$dF_t = \alpha_t F_t^\beta dW_t^1 \quad (1)$$

$$d\alpha_t = \nu \alpha_t dW_t^2 \quad (2)$$

where the two standard Brownian motions W_t^1 and W_t^2 are correlated such that $dW_t^1 dW_t^2 = \rho dt$.

The strength of the SABR model lies in the explicit financial interpretation of its four core parameters:

- **α (Alpha):** The initial volatility level. It dictates the overall height of the volatility curve, particularly At-The-Money (ATM).
- **β (Beta):** The structural parameter (often fixed by market convention, typically between 0 and 1). It determines the backbone of the model, dictating whether it behaves closer to a Normal/Bachelier model ($\beta = 0$) or a Lognormal/Black model ($\beta = 1$).
- **ρ (Rho):** The correlation between the asset price and its volatility. It controls the asymmetry or *skew* of the smile. A negative ρ (common in equity markets) produces a downward-sloping smile.
- **ν (Volvol):** The volatility of the volatility. It determines the convexity or *curvature* of the smile. Higher volvol makes the smile more pronounced (a deeper "U" shape).

1.2 The Machine Learning Paradigm: Introducing TabPFN

While neural networks have been widely used to learn pricing functions and speed up calibration, they generally require extensive, computationally expensive training phases (gradient descent over many epochs) for every new dataset.

TabPFN (Prior-Data Fitted Network for Tabular Data) represents a revolutionary step in machine learning. Based on the Transformer architecture, TabPFN does not require traditional training on the user's dataset. Instead, it utilizes **In-Context Learning (ICL)**. The model was pre-trained offline on millions of synthetic datasets generated from structural causal models.

When presented with a new problem, such as predicting implied volatility from financial parameters, TabPFN takes a small set of training examples as "context" and infers the mapping function in a single forward pass. The advantages for quantitative finance are significant:

1. **Instant Inference:** It completely bypasses the iterative optimization process, allowing for real-time market adaptation.
2. **Data Efficiency:** It achieves state-of-the-art performance even on very small datasets (e.g., sparse option chains with missing quotes).
3. **Bayesian Nature:** It inherently acts as an optimal Bayesian classifier/regressor, providing robust interpolations.

1.3 Bridging the Gap: Applying TabPFN to the SABR Model

The core objective of this project is to learn the complex, non-linear mapping between the market state/SABR parameters and the implied volatility:

$$(K, F, T, \alpha, \beta, \rho, \nu) \xrightarrow{\text{TabPFN}} \sigma_{\text{implied}} \quad (3)$$

By training an AI to understand the SABR dynamics, we aim to combine the theoretical soundness of stochastic calculus with the data-driven flexibility of Transformers. This allows us to "fill in the blanks" of an illiquid volatility surface instantly.

1.3.1 The Greek Problem and the Need for a Hybrid Architecture:

While TabPFN demonstrates exceptional accuracy in predicting the *value* of the volatility (the price), a critical issue arises in risk management. Traders do not just need prices; they require stable sensitivities, known as the **Greeks** (e.g., Delta $\Delta = \frac{\partial V}{\partial F}$, Vega $\mathcal{V} = \frac{\partial V}{\partial \alpha}$, Vanna, Volga).

Standard TabPFN models, like most purely data-driven neural networks, produce a highly erratic, noisy first derivative. A jagged gradient makes dynamic hedging impossible. To resolve this, this project goes beyond standard ML implementation. We propose a **Hybrid Stacking Architecture**:

- **The Prior (TabPFN):** Generates a highly accurate baseline prediction using in-context learning.
- **The Geometric Filter (MLP):** A deep Multi-Layer Perceptron that refines the TabPFN output.
- **Sobolev Differential Learning:** Using Automatic Differentiation (**JAX**), the model is trained not only on the SABR prices but heavily penalized on the *exact mathematical derivatives*.

2 SABR Engine Implementation (Steps 1 & 2)

2.1 Model Dynamics and Parameter Mapping (Step 1)

We implement the core SABR parameters—Alpha (α), Beta (β), Rho (ρ), and Volvol (ν)—to map the coupled stochastic differential equations of the forward price and its volatility. This numerical foundation allows the engine to translate any given market scenario into its corresponding implied volatility curve, effectively capturing the characteristic "smile" and skew observed in the options market.

2.2 Hagan's Asymptotic Approximation and Limitations (Step 2)

To compute the implied volatility efficiently, we utilize the widely adopted asymptotic expansion by Hagan et al. (2002). Using singular perturbation techniques, Hagan derived a closed-form approximation for the Black (lognormal) implied volatility. The derivation relies heavily on the assumption that the time to maturity T is relatively short and that the volatility of volatility ν is small, ensuring that the perturbation parameter $\epsilon = \nu^2 T$ remains bounded. The general form of the lognormal implied volatility σ_B is expressed as:

$$\sigma_B(K, F) \approx \frac{\alpha}{(FK)^{\frac{1-\beta}{2}} \left[1 + \frac{(1-\beta)^2}{24} \ln^2 \left(\frac{F}{K} \right) + \frac{(1-\beta)^4}{1920} \ln^4 \left(\frac{F}{K} \right) \right]} \cdot \left(\frac{z}{x(z)} \right) \cdot [1 + \Delta \cdot T]$$

where z represents a modified moneyness metric and $x(z)$ handles the correlation skew. This formula is incredibly fast to compute, making SABR the absolute industry standard for quoting swaptions and other interest rate derivatives. However, because it is merely a first-order asymptotic approximation, the formula is known to break down in extreme regimes. At very long maturities, highly negative interest rate environments, or deep Out-of-The-Money (OTM) strikes, it can produce severe anomalies, such as negative probability densities or arbitrage opportunities.

2.3 Singularity Management and Autodiff Implementation

A critical engineering challenge in coding Hagan's expansion is the strict numerical handling of the At-The-Money (ATM) singularity. When the strike approaches the forward price ($K \rightarrow F$), the term $z/x(z)$ evaluates to a 0/0 indeterminate form. We resolve this in the codebase by explicitly coding a limit-based ATM branch. This ensures strict mathematical continuity across the entire strike space, preventing gradient explosions (e.g., NaN values).

3 Synthetic Data Generation and Sobol Coverage (Step 3)

To train our models, we generate a large synthetic dataset based on the SABR framework. To ensure an optimal and unbiased exploration of the parameter space, we utilize a **Sobol sequence**. This method guarantees low-discrepancy coverage, preventing the clusters and voids typical of standard pseudo-random sampling.

3.1 Parameter Space and Scaling

All input features are sampled using Sobol sequences and mapped to a normalized interval of $[-1, 1]$ for neural network stability. Before normalization, the market scenarios strictly cover the following realistic bounds:

- **Model Parameters:** $\beta \in [0.25, 0.99]$, $\rho \in [-0.50, 0.50]$, and $\nu \in [0.15, 0.45]$.
- **Market Dynamics:** The Forward rate $F \in [1\%, 5\%]$, and the At-The-Money (ATM) normal volatility spans from 50 to 300 bps.
- **Strike Generation:** To ensure financial relevance, strikes (K) are not generated purely at random but are anchored to the Forward via a moneyness factor (K/F) ranging from 75% to 150%.

3.2 Dataset Analysis and Structural Integrity

As illustrated in Figure 1, the pairplot analysis confirms the quality of our data distribution.

The diagonal histograms show perfectly uniform plateaus for β , ρ , and ν , validating the Sobol uniformity. Notably, the distribution of **Alpha** (α) is concentrated near zero; this is not a sampling bias but a direct result of the non-linear inversion of Hagan’s formula from the uniform ATM volatility space.

Bivariate plots reveal the "cone" structure between F and K , reflecting the strict mathematical coupling defined by our moneyness bounds. The near-zero Pearson correlations between the other model parameters confirm their independence, ensuring that the subsequent AI models learn the specific sensitivity of each risk factor without multicollinearity bias.




Figure 1: SABR Parameter Space Analysis. Diagonal histograms show marginal distributions, while bivariate scatter plots reveal the structural dependencies.

4 Baseline Training: Standard TabPFN (Step 4)

With the dataset generated, we apply the **TabPFN** model. TabPFN is a transformer-based foundation model that performs Bayesian inference on tabular data through a single forward pass.

4.1 Pre-processing and Scaling

To optimize the learning process, we implement a sophisticated pre-processing pipeline:

- **Log-Transformations:** We compute the *log-moneyness* $\ln(K/F)$, which is a more natural coordinate system for volatility models.
- **MinMaxScaler:** All input features are scaled between $[-1, 1]$ to ensure numerical stability within the Transformer’s attention mechanisms.
- **Output Normalization:** The target volatility is scaled between $[0, 1]$.

4.2 Initial Performance and The "Gradient Challenge"




Figure 2: TabPFNP vs SABR on volatility

The initial results of the standard TabPFN are impressive in terms of price accuracy. As illustrated in Figure 2, the model achieves a remarkably low Mean Absolute Error (MAE) of 2.85×10^{-5} on the implied volatility values, demonstrating a near-perfect diagonal alignment between the model’s predictions and the SABR ground truth.

5 The Derivatives Issue: Baseline AI Evaluation

Using our normalized Sobol dataset, we evaluate the baseline performance of the standard **TabPFN** model. While the model demonstrates an excellent capacity to predict implied volatility, a geometric analysis reveals a critical structural flaw.

5.1 JAX as the Analytical Ground Truth

To accurately measure this geometric failure, we require a flawless benchmark. We utilize our fully differentiable **JAX** pricing engine to serve as the absolute mathematical reference. By leveraging JAX’s automatic differentiation, we can instantly compute the exact.

5.2 Comparative Analysis: The Gradient Challenge

With the JAX ground truth established, we extracted the empirical Skew from the TabPFN predictions via finite differences. To isolate the shape and geometric instability, we applied a Z-Score alignment to both trajectories.




Figure 3: TabPFN vs JAX skew

As illustrated in Figure 3, the AI’s numerical derivative oscillates wildly around the true mathematical trajectory. In a real-world trading environment, these jagged variations would translate into unstable, erratic, and highly costly hedging signals. This specific failure—high accuracy in value but chaos in derivatives—serves as the definitive proof that standard AI models cannot be deployed directly in trading systems, thereby motivating the introduction of Sobolev learning in the subsequent steps.

6 Sobolev Loss Architecture: Learning the Surface Geometry (Step 5)

To resolve the geometric instability identified in Step 4, we transition from a standard point-wise regression to a **Sobolev-informed learning** framework. The core idea is to force the neural network to minimize not only the error on the volatility values but also the discrepancy between its internal gradients and the analytical Greeks provided by JAX.

6.1 Mathematical Formulation

We implement a custom `DerivativeLoss` module in PyTorch. The total loss function $\mathcal{L}_{total}$ is defined as a weighted sum of the volatility error and the average sensitivity error:

$$\mathcal{L}_{total} = w_v \cdot \text{MAE}(\hat{V}, V) + w_d \cdot \frac{1}{N} \sum_{i=1}^N \text{MAE}(\nabla \hat{V}_i, \nabla V_i) \quad (4)$$

where:

- $\hat{V}$ and V are the predicted and true implied volatilities.
- $\nabla \hat{V}_i$ represents the i -th predicted sensitivity (e.g., $\partial V / \partial K$, $\partial V / \partial \rho$).
- w_v and w_d are hyper-parameters controlling the trade-off between price accuracy and surface smoothness.

6.2 Implementation and Monitoring

A key feature of our implementation is the generation of a **Loss Breakdown** dictionary during each training step. Rather than monitoring a single scalar, we track the individual Mean Absolute Error (MAE) for every Greek (Δ, ν, ρ , etc.).

This granular monitoring is crucial for two reasons:

1. **Convergence Analysis:** It allows us to identify which financial sensitivities are more difficult for the model to learn (e.g., the Skew near the ATM singularity).
2. **Hyper-parameter Tuning:** It provides the necessary feedback to adjust the derivative weight w_d , ensuring that the gradient regularization is strong enough to suppress numerical noise without sacrificing the global accuracy of the volatility surface.

7 Data Generation Strategy: Hybrid Mesh Refinement (Step 5)

While a standard Sobol sequence provides excellent uniform coverage of the parameter hypercube, financial models such as SABR exhibit high non-linearities that a naive uniform grid cannot fully capture. The most critical and numerically sensitive region is the At-The-Money (ATM) singularity, where the strike equals the forward ($K \approx F$). In this area, the volatility curvature and the resulting gradients (the Skew and Smile) reach their peak, requiring a much higher informational granularity.

To ensure the stability of our neural solver, we implemented a **Hybrid Mesh Refinement** strategy. This approach moves away from basic uniform sampling to a dual-layer distribution specifically designed to "anchor" the dataset around the model's mathematical core while maintaining a robust boundary coverage.

7.1 Methodology: Sobol Foundation and Gaussian ATM Refinement

The Hybrid Mesh algorithm operates by merging two distinct sampling philosophies to create a heterogeneous density of points:

1. **Global Mesh (Sobol Sequence):** First, we generate a low-discrepancy Sobol sequence to cover the entire parametric space. Unlike pseudo-random sampling, this ensures that the edges of the model (extreme OTM/ITM strikes and long maturities) are well-represented, preventing the network from *data starvation* at the boundaries.
2. **Local Mesh Refinement (Gaussian ATM Layer):** To resolve the ATM singularity, we inject a hyper-dense layer of points centered at the moneyness origin. We perform this **mesh refinement** by drawing log-moneyness values ($\ln(K/F)$) from a tight Gaussian distribution ($\mu = 0, \sigma = 0.10$). This forces a significant portion of the training data (50%) to be clustered exactly where $K/F \approx 1$.

As illustrated in Figure 4, this **refinement strategy** (bottom panel) effectively resolves the structural limitations of standard uniform sampling (top panel). By targeted reallocation of data density, we ensure that the peak of the sample allocation is strictly anchored at the $K/F = 1.0$ singularity. This "structural anchoring" via mesh refinement guarantees that the neural network receives maximum informational granularity in the regions most prone to interpolation errors. Ultimately, this targeted data engineering is what secures the geometric stability and the smoothness of the Greeks across the entire volatility surface.




Figure 4: Comparison: Uniform Sobol vs. Hybrid Mesh Refinement Distribution

8 Architecture Search and Hyperparameter Optimization (Step 7)

To identify the optimal neural network topology for learning the SABR volatility surface and its associated Greeks, we performed an automated hyperparameter search using **Ray Tune** (leveraging the Optuna search algorithm and the ASHA scheduler for early stopping).

8.1 Experimental Setup and Search Space

The models were trained and evaluated exclusively on the high-density **Adaptive Mesh Refinement (AMR) dataset** generated in Step 5. The primary optimization objective was to minimize the **Validation Sobolev Loss**, ensuring the model accurately predicts both the volatility values and the 6 financial sensitivities.

Each trial was trained for a maximum of **50 epochs**. The search space was defined as follows:

- **Base Architectures:** A comparison between standard Deep MLPs and Tabular Transformers (varying $d_{\text{model}} \in \{128, 256, 512\}$, $n_{\text{head}} \in \{4, 8\}$, and depth).
- **Regression Head:** Dense layers defined by hidden dimensions $\in \{(512, 256, 128), (256, 128, 64)\}$ to map the latent representation to the 7 continuous outputs.
- **Activations:** Strictly twice-differentiable functions (Swish, Mish, GELU, SELU) to preserve uninterrupted gradient flow for the JAX derivatives.

- **Training Dynamics:** Batch size, learning rate, dropout, and weight decay (via Adam / AdamW).
- **Sobolev Regularization:** Autonomous tuning of the custom loss parameters (Value weight $\in [0.5, 1.5]$ and Derivative weight $\in [0.1, 1.0]$).

8.2 Optimal Configuration

The optimization process evaluated multiple topologies and ultimately converged on a high-capacity **Multi-Layer Perceptron (MLP)** architecture integrated within the stacking framework. As detailed in Table 1, the winning configuration utilizes a hybrid depth strategy, featuring a latent representation dimension of $d_{\text{model}} = 256$ with 6 transformer-inspired layers for feature extraction, followed by a specialized deep [512, 256, 128] MLP head. Notably, the engine selected the **GELU** (Gaussian Error Linear Unit) activation function; being smooth and non-monotonic, it ensures a stable and uninterrupted gradient flow during the backpropagation of the JAX sensitivities, which is vital for capturing the SABR model’s curvature.

Crucially, the search mathematically validated the necessity of our Sobolev constraint. The optimization engine dynamically balanced the learning objectives, assigning a weight of 6.53×10^{-1} to the volatility value and 1.32×10^{-1} to the derivative loss. This specific ratio ensures that the network does not merely "memorize" price points but strictly adheres to the geometric laws of the SABR manifold, guaranteeing stable risk sensitivities even at the ATM singularity.

Hyperparameters	Values
Model type	MLP
Activation	GELU
d_{model}	256
Number of heads (n_{head})	8
Number of layers	6
Feedforward dimension	2048
Hidden dimensions (MLP Head)	[512, 256, 128]
Batch size	128
Learning rate	9.32×10^{-3}
Dropout	9.81×10^{-2}
Optimizer	Adam
Weight decay	5.67×10^{-5}
Value loss weight (λ_v)	6.53×10^{-1}
Derivative loss weight (λ_d)	1.32×10^{-1}

Table 1: Optimal Hyperparameters of the SABR Hybrid Model determined by Ray Tune search.

9 Stacking (Step 8)

The final step of our methodology consists of assembling the previously developed modules (the JAX engine, TabPFN, and the optimized MLP) into a cohesive **Stacking** architecture. This hybrid approach is designed to overcome the geometric limitations of foundation models while preserving their robust predictive power.

9.1 Stacking Architecture and Feature Augmentation

Rather than training a deep learning model in isolation, we deploy **TabPFN** as a pricing oracle. For each configuration in the AMR dataset, TabPFN generates an initial implied volatility estimate. This baseline prediction is then injected as an **augmented feature** (Rich Feature) alongside the original 8 SABR parameters.

Consequently, the final model (the **MLP Stacker**) processes a 9-dimensional input vector. This configuration allows the neural network to bypass the arduous task of learning option pricing from scratch. Instead, it dedicates its representational capacity to adjusting and mathematically smoothing the highly accurate, yet geometrically unstable, prediction provided by TabPFN.




Figure 5: Architecture of the Hybrid Reconstruction Pipeline. The data flow combines the pricing expertise of TabPFN with the geometric regularization of the MLP, strictly supervised by the exact gradients of the JAX engine.

9.2 Sobolev Loss Regularization and JAX Ground Truth

As illustrated in Figure 5, the training phase is governed by a co-optimized **Sobolev Loss**. This dual-supervision framework penalizes two distinct sources of error:

1. **Value Error:** The model minimizes the absolute deviation between the predicted volatility and the JAX ground truth.
2. **Derivative Error (Greeks):** Simultaneously, the model's internal gradients (computed via automatic differentiation during backpropagation) are compared against the exact analytical sensitivities generated by the **JAX** engine (Step 2).

This strict gradient constraint forces the network to learn not merely the "points" on the volatility surface, but its true **mathematical curvature**. By acting as a regularizing filter over the TabPFN predictions, the MLP guarantees a perfectly smooth volatility surface and highly stable Greeks (Δ, ν, ρ , etc.). This mechanism conclusively resolves the "Gradient Challenge" identified in the earlier stages of our research.

10 Finetuning on Causal Graph

need to be improved

11 Results and Geometric Validation

The initial discrepancy observed during the 3D validation phase was ultimately diagnosed not as a failure of the neural architecture, but as a fundamental paradigm mismatch in the evaluation framework. The apparent collapse of the predicted surface was caused by an inadvertent comparison between neural predictions—strictly trained on lognormal (Black) implied volatilities—and a normal (Bachelier) mathematical baseline. Normal volatilities inherently yield microscopically small gradients compared to their lognormal counterparts. By realigning the evaluation script to strictly operate within the true lognormal SABR space, and by dynamically centering the test scenario (e.g., $F = 0.04$, $\alpha \approx 0.22$) within the core distribution of the standardized training data, the artificial out-of-distribution (OOD) extrapolation anomalies were entirely mitigated.

The perfect superimposition of the surfaces is the direct validation of the two major architectural breakthroughs introduced in this study: the Hybrid Mesh generation and the Adaptive Sobolev Loss. First, the Hybrid Mesh sampling strategy proved critical. By combining a global Sobol sequence covering the entire parameter hypercube with an At-The-Money (ATM) localized Gaussian refinement, the training dataset guaranteed sufficient data density at the deep In-The-Money (ITM) and Out-Of-The-Money (OTM) boundaries. This explicitly solved the "data starvation" phenomenon that previously caused the Multi-Layer Perceptron (MLP) to diverge catastrophically at the edges of the volatility surface. The network is no longer forced into blind extrapolation; instead, it interpolates smoothly across a well-defined topological space.

Secondly, the structural integrity of the predicted surface highlights the power of the *Adaptive Sobolev Loss* function based on homoscedastic uncertainty weighting. Rather than relying on static, heuristically determined penalty weights (found via hyperparameter optimization), the network autonomously learned the optimal balance during training. This allowed the optimizer to dynamically weigh the Mean Absolute Error (MAE) of the volatility values against the highly sensitive gradients (the Greeks), particularly the volatility skew ($\partial V / \partial K$). Consequently, the network prioritizes the fundamental shape of the surface just as much as the absolute price levels.

As clearly illustrated in Figure 6, the AI prediction (orange surface) natively hugs the exact mathematical truth of the Hagan Lognormal SABR formulation (translucent green surface) across the entire Forward-Strike (F, K) domain. The model successfully captures the highly nonlinear dynamics of the volatility skew at its true scale (yielding structurally accurate values in the vicinity of $z \approx -18$), completely eliminating the need for any artificial statistical realignment, Z-score matching, or ad-hoc post-processing.




Figure 6: AI prediction (orange surface) perfectly overlaps the true skew (green surface)

This absolute structural convergence validates the TabPFN-augmented stacking architecture. By leveraging the fast, in-context learning capabilities of Case-Based Reasoning (TabPFN) as an oracle, and a deep MLP stacker penalized by differential equations for continuous smoothing, the pipeline acts as a highly robust, data-efficient solver. It demonstrates that deep learning models can indeed preserve the fundamental, arbitrage-free geometries required in quantitative finance, effectively resolving the "Gradient Challenge" and offering a viable pathway for high-fidelity, real-time risk management applications.